

```

import numpy as np
import pandas as pd
import sklearn
from sklearn.datasets import load_boston
boston = load_boston()
df = pd.DataFrame(boston.data, columns=boston.feature_names)
df.head()
df['MEDV'] = boston.target
df.head()
df.isnull()
df.isnull().sum()
from sklearn.model_selection import train_test_split
X_train, X_test, Y_train, Y_test = train_test_split(df, df['MEDV'], test_size = 0.15, random_state= 5)
lin_model = LinearRegression()
lin_model.fit(X_train, Y_train)
lin_model.predict(X_train)
(np.sqrt(mean_squared_error(Y_train, lin_model.predict(X_train))))
"The model performance for training set"
"RMSE is {}".format(rmse)
"\n")
lin_model.predict(X_test)
(np.sqrt(mean_squared_error(Y_test, lin_model.predict(X_test))))
"The model performance for testing set"
"RMSE is {}".format(rmse)

```

495	23.1
206	24.4
355	20.6
Name: MEDV, Length: 430, dtype: float64	
226	37.6
292	27.9
90	22.6
373	13.8
273	35.2
...	
311	22.1
49	19.4
409	27.5
407	27.9
353	30.1
Name: MEDV, Length: 76, dtype: float64	
The model performance for training set	
RMSE is 4.710901797319796	

```

The model performance for testing set
RMSE is 4.687543527902972
/usr/local/lib/python3.7/dist-packages/sklearn/utils/deprecation.py:87: FutureWarning:

```

The Boston housing prices dataset has an ethical problem. You can refer to the documentation of this function for further details.

The scikit-learn maintainers therefore strongly discourage the use of this dataset unless the purpose of the code is to study and educate about ethical issues in data science and machine learning.

In this special case, you can fetch the dataset from the original source::

```
import pandas as pd
import numpy as np

data_url = "http://lib.stat.cmu.edu/datasets/boston"
raw_df = pd.read_csv(data_url, sep="\s+", skiprows=22, header=None)
data = np.hstack([raw_df.values[::2, :], raw_df.values[1::2, :2]])
target = raw_df.values[1::2, 2]
```

Alternative datasets include the California housing dataset (i.e. :func:`~sklearn.datasets.fetch\_california\_housing`) and the Ames housing dataset. You can load the datasets as follows::

```
from sklearn.datasets import fetch_california_housing
housing = fetch_california_housing()
```

for the California housing dataset and::

```
from sklearn.datasets import fetch_openml
housing = fetch_openml(name="house_prices", as_frame=True)
```

for the Ames housing dataset.

```
warnings.warn(msg, category=FutureWarning)
```